The Azubi Method

Introduction and Overview

Structured Naïveté as a Review Mechanism for AI Agents

Andreas Schmidt

andreas-schmidt-arts.de

March 2026

What This Paper Is

This paper describes a method for quality assurance of results produced by AI agents. Not through more expertise, but through the structured absence of expertise. Not through better answers, but through naive questions.

The method is called the Azubi Method and formalizes a principle that has been known in the human workplace for generations: The new employee on their first day asks the question that everyone else considers too obvious. And that very question uncovers the deepest problem.

The Azubi Method was developed in software quality assurance and validated over more than 20 audit rounds across four products. But its scope is not limited to code. It can be applied to any result produced by an expert or an expert system: assessments, analyses, expert opinions, audit reports, strategy documents, security audits, risk evaluations. Anywhere an LLM produces a result and that result should be questioned.

This paper is a practice report with research connections: hypotheses are marked as hypotheses, limitations are named, sources are verifiable. It is part of a larger paper series ("Multi-Agent QA Framework with the Azubi Method," seven parts), from which it extracts the core method as a standalone document.

Who This Paper Is For

For developers and teams who use AI agents for review and evaluation and want to know why their results still have gaps.

For companies that develop language models. The mechanisms described here could flow directly into future models. A model that automatically poses a context-free counter-question after a review would no longer need an external Azubi.

For research institutions and universities. This paper formulates open questions suited for scientific investigation: Can the effect of structured naïveté be measured empirically? How does the method behave in controlled multi-agent scenarios?

Contents

Part 1: Why Expertise Creates Blind Spots. The ELIZA effect, sycophancy, operational blindness in AI agents. Three error classes and the attention threshold. Why existing review methods fail to cover a specific class of errors.

Part 2: The Azubi Method. The core: context-free naïveté as a review mechanism. Five properties, two question types, practical results. Why it works technically: breaking reward hacking. The expert-Azubi dual play. Anti-patterns.

Part 3: Application, Limitations, and Future. General applicability beyond code. Related approaches. Honestly stated limitations. Three levers for LLM developers. Conclusion, references, about the author.